

Introduction

Welcome to **CS188 - Data Science Fundamentals!** This course is designed to equip you with the tools and experiences necessary to start you off on a life-long exploration of datascience. We do not assume a prerequisite knowledge or experience in order to take the course.

For this first project we will introduce you to the end-to-end process of doing a datascience project. Our goals for this project are to:

1. Familiarize you with the development environment for doing datascience
2. Get you comfortable with the python coding required to do datascience
3. Provide you with an sample end-to-end project to help you visualize the steps needed to complete a project on your own
4. Ask you to recreate a similar project on a separate dataset

In this project you will work through an example project end to end. Many of the concepts you will encounter will be unclear to you. That is OK! The course is designed to teach you these concepts in further detail. For now our focus is simply on having you replicate the code successfully and seeing a project through from start to finish.

Here are the main steps:

1. Get the data
2. Visualize the data for insights
3. Preprocess the data for your machine learning algorithm
4. Select a model and train
5. Does it meet the requirements? Fine tune the model



Working with Real Data

It is best to experiment with real-data as opposed to artificial datasets.

There are many different open datasets depending on the type of problems you might be interested in!

Here are a few data repositories you could check out:

- [UCI Datasets](#)
- [Kaggle Datasets](#)
- [AWS Datasets](#)

Submission Instructions

When you have completed this assignment please save the notebook as a PDF file and submit the assignment via Gradescope

Example Datascience Exercise

Below we will run through an California Housing example collected from the 1990's.

Setup

```
In [90]: import sys
assert sys.version_info >= (3, 5) # python>=3.5
import sklearn
assert sklearn.__version__ >= "0.20" # sklearn >= 0.20

import numpy as np #numerical package in python
import os
%matplotlib inline
import matplotlib.pyplot as plt #plotting package

# to make this notebook's output identical at every run
np.random.seed(42)

#matplotlib magic for inline figures
%matplotlib inline
import matplotlib # plotting library
import matplotlib.pyplot as plt

# Where to save the figures
ROOT_DIR = "."
IMAGES_PATH = os.path.join(ROOT_DIR, "images")
os.makedirs(IMAGES_PATH, exist_ok=True)

def save_fig(fig_name, tight_layout=True, fig_extension="png", resolution=300):
    """
    plt.savefig wrapper. refer to
    https://matplotlib.org/3.1.1/api/_as_gen/matplotlib.pyplot.savefig.html

    Args:
        fig_name (str): name of the figure
        tight_layout (bool): adjust subplot to fit in the figure area
        fig_extension (str): file format to save the figure in
        resolution (int): figure resolution
    """
    path = os.path.join(IMAGES_PATH, fig_name + "." + fig_extension)
    print("Saving figure", fig_name)
    if tight_layout:
        plt.tight_layout()
    plt.savefig(path, format=fig_extension, dpi=resolution)
```

```
In [91]: import os
import tarfile
import urllib
DATASET_PATH = os.path.join("datasets", "housing")
```

Step 1. Getting the data

Intro to Data Exploration Using Pandas

In this section we will load the dataset, and visualize different features using different types of plots.

Packages we will use:

- **Pandas**: is a fast, flexible and expressive data structure widely used for tabular and multidimensional datasets.
- **Matplotlib**: is a 2d python plotting library which you can use to create quality figures (you can plot almost anything if you're willing to code it out!)
 - other plotting libraries: [seaborn](#), [ggplot2](#)

```
In [92]: import pandas as pd

def load_housing_data(housing_path):
    """
        loads housing.csv dataset stored

        Args:
            housing_path (str): path to folder containing housing dataset

        Returns:
            pd.DataFrame
    """
    csv_path = os.path.join(housing_path, "housing.csv")
    return pd.read_csv(csv_path)
```

```
In [93]: pd.DataFrame
```

```
Out[93]: pandas.core.frame.DataFrame
```

```
In [94]: housing = load_housing_data(DATASET_PATH) # we load the pandas dataframe
housing.head() # show the first few elements of the dataframe
               # typically this is the first thing you do
               # to see how the dataframe looks like
```

```
Out[94]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households
0	-122.23	37.88	41.0	880.0	129.0	322.0	126.0
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0	1138.0
2	-122.24	37.85	52.0	1467.0	190.0	496.0	177.0
3	-122.25	37.85	52.0	1274.0	235.0	558.0	219.0
4	-122.25	37.85	52.0	1627.0	280.0	565.0	259.0

A dataset may have different types of features

- real valued
- Discrete (integers)
- categorical (strings)

The two categorical features are essentially the same as you can always map a categorical string/character to an integer.

In the dataset example, all our features are real valued floats, except ocean proximity which is categorical.

```
In [95]: # to see a concise summary of data types, null values, and counts  
# use the info() method on the dataframe  
housing.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 20640 entries, 0 to 20639  
Data columns (total 10 columns):  
#   Column                Non-Null Count  Dtype  
---  ---  
0   longitude              20640 non-null  float64  
1   latitude               20640 non-null  float64  
2   housing_median_age     20640 non-null  float64  
3   total_rooms            20640 non-null  float64  
4   total_bedrooms         20433 non-null  float64  
5   population             20640 non-null  float64  
6   households             20640 non-null  float64  
7   median_income          20640 non-null  float64  
8   median_house_value     20640 non-null  float64  
9   ocean_proximity        20640 non-null  object  
dtypes: float64(9), object(1)  
memory usage: 1.6+ MB
```

```
In [96]: # you can access individual columns similarly  
# to accessing elements in a python dict  
housing["ocean_proximity"].head() # added head() to avoid printing many columns.
```

```
Out[96]: 0    NEAR BAY  
1    NEAR BAY  
2    NEAR BAY  
3    NEAR BAY  
4    NEAR BAY  
Name: ocean_proximity, dtype: object
```

```
In [97]: # to access a particular row we can use iloc  
housing.iloc[1]
```

```
Out[97]: longitude          -122.22  
latitude              37.86  
housing_median_age      21.0  
total_rooms            7099.0  
total_bedrooms          1106.0  
population             2401.0  
households             1138.0  
median_income           8.3014  
median_house_value     358500.0  
ocean_proximity        NEAR BAY  
Name: 1, dtype: object
```

```
In [98]: # one other function that might be useful is  
# value_counts(), which counts the number of occurrences  
# for categorical features  
housing["ocean_proximity"].value_counts()
```

```
Out[98]: <1H OCEAN      9136
INLAND          6551
NEAR OCEAN      2658
NEAR BAY        2290
ISLAND           5
Name: ocean_proximity, dtype: int64
```

```
In [99]: # The describe function compiles your typical statistics for each
# column
housing.describe()
```

```
Out[99]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population
count	20640.000000	20640.000000	20640.000000	20640.000000	20433.000000	20640
mean	-119.569704	35.631861	28.639486	2635.763081	537.870553	1425
std	2.003532	2.135952	12.585558	2181.615252	421.385070	1132
min	-124.350000	32.540000	1.000000	2.000000	1.000000	3
25%	-121.800000	33.930000	18.000000	1447.750000	296.000000	787
50%	-118.490000	34.260000	29.000000	2127.000000	435.000000	1166
75%	-118.010000	37.710000	37.000000	3148.000000	647.000000	1725
max	-114.310000	41.950000	52.000000	39320.000000	6445.000000	35682

If you want to learn about different ways of accessing elements or other functions it's useful to check out the getting started section [here](#)

Step 2. Visualizing the data

Let's start visualizing the dataset

```
In [100... # We can draw a histogram for each of the dataframes features
# using the hist function
housing.hist(bins=50, figsize=(20,15))
# save_fig("attribute_histogram_plots")
plt.show() # pandas internally uses matplotlib, and to display all the figures
# the show() function must be called
```




```
labels=[1, 2, 3, 4, 5])

housing["income_cat"].value_counts()
```

```
Out[102...] 3    7236
            2    6581
            4    3639
            5    2362
            1     822
            Name: income_cat, dtype: int64
```

```
In [103...] housing["income_cat"].hist()
```


```
Out[103...] <matplotlib.axes._subplots.AxesSubplot at 0x7fb38bbb3910>
```



Next let's visualize the household incomes based on latitude & longitude coordinates

```
In [104...] ## here's a not so interesting way of plotting it
housing.plot(kind="scatter", x="longitude", y="latitude")
save_fig("bad_visualization_plot")
```


Saving figure bad_visualization_plot



```
In [105...] # we can make it look a bit nicer by using the alpha parameter,
```

```
# it simply plots less dense areas lighter.
housing.plot(kind="scatter", x="longitude", y="latitude", alpha=0.1)
save_fig("better_visualization_plot")
```

Saving figure better_visualization_plot



In [106...

```
# A more interesting plot is to color code (heatmap) the dots
# based on income. The code below achieves this

# load an image of california
images_path = os.path.join('.', 'images')
os.makedirs(images_path, exist_ok=True)
filename = "california.png"

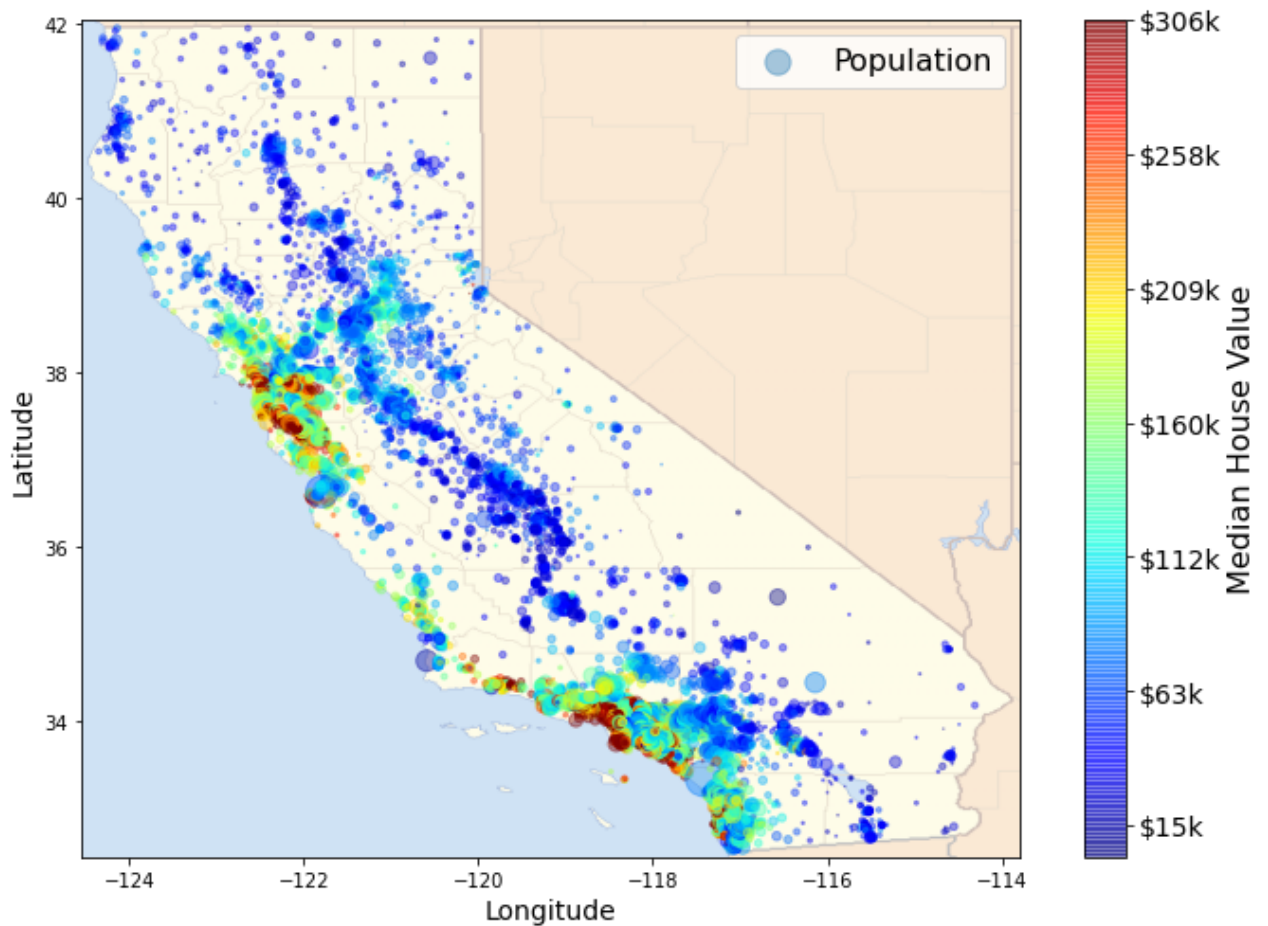
import matplotlib.image as mpimg
california_img=mpimg.imread(os.path.join(images_path, filename))
ax = housing.plot(kind="scatter", x="longitude", y="latitude", figsize=(10,7),
                  s=housing['population']/100, label="Population",
                  c="median_house_value", cmap=plt.get_cmap("jet"),
                  colorbar=False, alpha=0.4,
                  )

# overlay the califronia map on the plotted scatter plot
# note: plt.imshow still refers to the most recent figure
# that hasn't been plotted yet.
plt.imshow(california_img, extent=[-124.55, -113.80, 32.45, 42.05], alpha=0.5,
           cmap=plt.get_cmap("jet"))
plt.ylabel("Latitude", fontsize=14)
plt.xlabel("Longitude", fontsize=14)

# setting up heatmap colors based on median_house_value feature
prices = housing["median_house_value"]
tick_values = np.linspace(prices.min(), prices.max(), 11)
cb = plt.colorbar()
cb.ax.set_yticklabels(["$%dk"%(round(v/1000)) for v in tick_values], fontsize=14)
cb.set_label('Median House Value', fontsize=16)

plt.legend(fontsize=16)
save_fig("california_housing_prices_plot")
plt.show()
```

Saving figure california_housing_prices_plot



Not surprisingly, we can see that the most expensive houses are concentrated around the San Francisco/Los Angeles areas.

Up until now we have only visualized feature histograms and basic statistics.

When developing machine learning models the predictiveness of a feature for a particular target of interest is what's important.

It may be that only a few features are useful for the target at hand, or features may need to be augmented by applying certain transformations.

None the less we can explore this using correlation matrices. If you need to brush up on correlation take a look [here](#).

```
In [107... corr_matrix = housing.corr() # compute the correlation matrix
```

```
In [108... # for example if the target is "median_house_value", most correlated features ca
# which happens to be "median_income". This also intuitively makes sense.
corr_matrix["median_house_value"].sort_values(ascending=False)
```

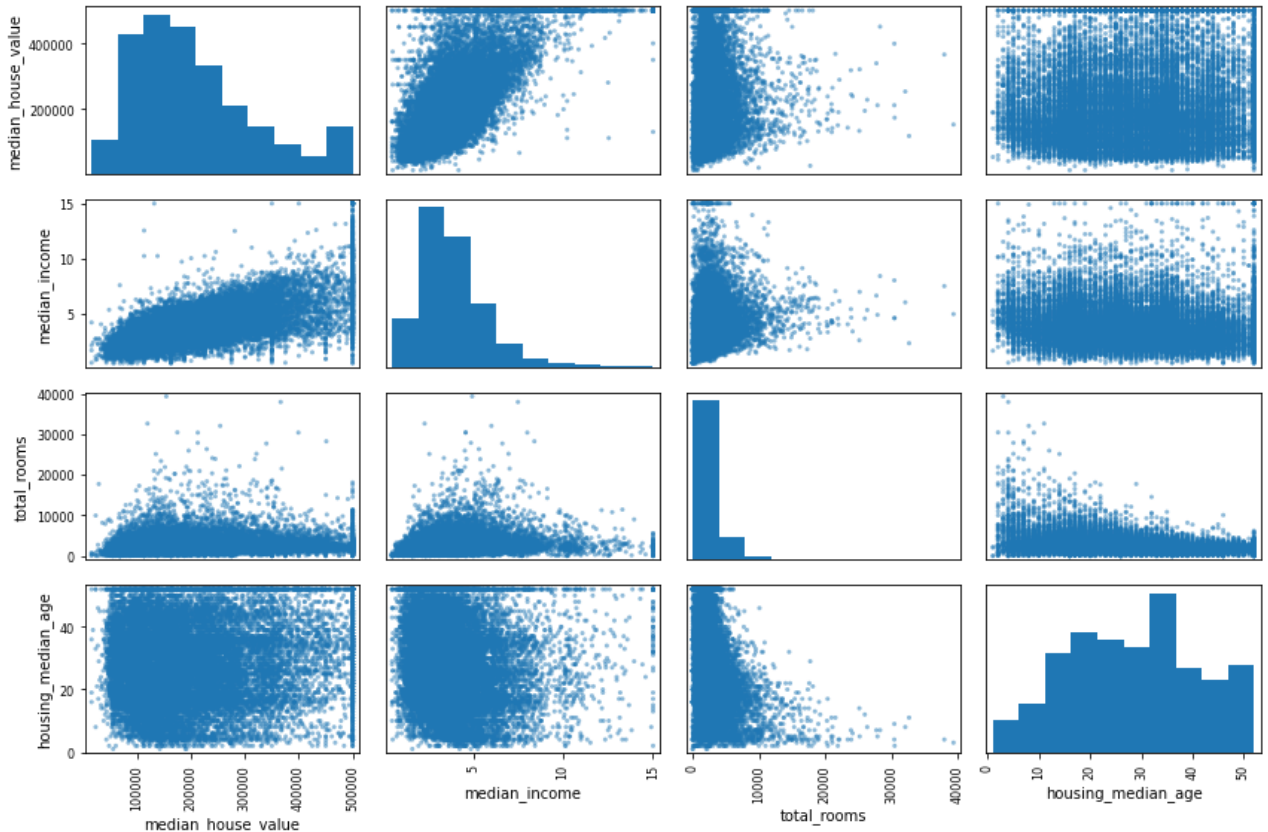
```
Out[108... median_house_value    1.000000
median_income        0.688075
total_rooms          0.134153
housing_median_age    0.105623
households           0.065843
total_bedrooms       0.049686
```

```
population          -0.024650
longitude           -0.045967
latitude            -0.144160
Name: median_house_value, dtype: float64
```

In [109...

```
# the correlation matrix for different attributes/features can also be plotted
# some features may show a positive correlation/negative correlation or
# it may turn out to be completely random!
from pandas.plotting import scatter_matrix
attributes = ["median_house_value", "median_income", "total_rooms",
             "housing_median_age"]
scatter_matrix(housing[attributes], figsize=(12, 8))
save_fig("scatter_matrix_plot")
```


Saving figure scatter_matrix_plot



In [110...

```
# median income vs median house value plot plot 2 in the first row of top figure
housing.plot(kind="scatter", x="median_income", y="median_house_value",
               alpha=0.1)
plt.axis([0, 16, 0, 550000])
save_fig("income_vs_house_value_scatterplot")
```

Saving figure income_vs_house_value_scatterplot



Augmenting Features

New features can be created by combining different columns from our data set.


- $\text{rooms_per_household} = \text{total_rooms} / \text{households}$
- $\text{bedrooms_per_room} = \text{total_bedrooms} / \text{total_rooms}$
- etc.

```
In [111... housing["rooms_per_household"] = housing["total_rooms"]/housing["households"]
housing["bedrooms_per_room"] = housing["total_bedrooms"]/housing["total_rooms"]
housing["population_per_household"] = housing["population"]/housing["households"]
```


```
In [112... # obtain new correlations
corr_matrix = housing.corr()
corr_matrix["median_house_value"].sort_values(ascending=False)
```

```
Out[112... median_house_value      1.000000
median_income          0.688075
rooms_per_household    0.151948
total_rooms            0.134153
housing_median_age     0.105623
households             0.065843
total_bedrooms         0.049686
population_per_household -0.023737
population             -0.024650
longitude              -0.045967
latitude               -0.144160
bedrooms_per_room      -0.255880
Name: median_house_value, dtype: float64
```

```
In [113... housing.plot(kind="scatter", x="rooms_per_household", y="median_house_value",
                    alpha=0.2)
plt.axis([0, 5, 0, 520000])
plt.show()
```



```
In [114]: housing.plot(kind="scatter", x="median_income", y="median_house_value",
alpha=0.2)
plt.axis([0, 5, 0, 520000])
plt.show()
```



```
In [115]: housing.describe()
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population
count	20640.000000	20640.000000	20640.000000	20640.000000	20433.000000	20640
mean	-119.569704	35.631861	28.639486	2635.763081	537.870553	1425
std	2.003532	2.135952	12.585558	2181.615252	421.385070	1132
min	-124.350000	32.540000	1.000000	2.000000	1.000000	3
25%	-121.800000	33.930000	18.000000	1447.750000	296.000000	787
50%	-118.490000	34.260000	29.000000	2127.000000	435.000000	1166
75%	-118.010000	37.710000	37.000000	3148.000000	647.000000	1725
max	-114.310000	41.950000	52.000000	39320.000000	6445.000000	35682

Step 3. Preprocess the data for your machine learning algorithm

Once we've visualized the data, and have a certain understanding of how the data looks like. It's time to clean!

Most of your time will be spent on this step, although the datasets used in this project are relatively nice and clean... in the real world it could get real dirty.

After having cleaned your dataset you're aiming for:

- train set
- test set

In some cases you might also have a validation set as well for tuning hyperparameters (don't worry if you're not familiar with this term..)

In supervised learning setting your train set and test set should contain (**feature**, **target**) tuples.

- **feature**: is the input to your model
- **target**: is the ground truth label
 - when target is categorical the task is a classification task
 - when target is floating point the task is a regression task

We will make use of [scikit-learn](#) python package for preprocessing.

Scikit learn is pretty well documented and if you get confused at any point simply look up the function/object!

Dealing With Incomplete Data

```
In [116... # have you noticed when looking at the dataframe summary certain rows  
# contained null values? we can't just leave them as nulls and expect our  
# model to handle them for us so we'll have to devise a method for dealing with  
sample_incomplete_rows = housing[housing.isnull().any(axis=1)].head()  
sample_incomplete_rows
```

```
Out[116...      longitude  latitude  housing_median_age  total_rooms  total_bedrooms  population  households  
290      -122.16    37.77             47.0      1256.0             NaN        570.0         21  
341      -122.17    37.75             38.0       992.0             NaN        732.0         25  
538      -122.28    37.78             29.0     5154.0             NaN       3741.0        127  
563      -122.24    37.75             45.0       891.0             NaN        384.0         14  
696      -122.10    37.69             41.0       746.0             NaN        387.0         16
```

```
In [117... sample_incomplete_rows.dropna(subset=["total_bedrooms"]) # option 1: simply d
```

```
Out[117...      longitude  latitude  housing_median_age  total_rooms  total_bedrooms  population  households
```

```
In [118... sample_incomplete_rows.drop("total_bedrooms", axis=1) # option 2: drop the
```

```
Out[118...
      longitude  latitude  housing_median_age  total_rooms  population  households  median_inco
290    -122.16    37.77             47.0      1256.0      570.0        218.0        4.3
341    -122.17    37.75             38.0       992.0      732.0        259.0        1.6
538    -122.28    37.78             29.0     5154.0     3741.0       1273.0        2.5
563    -122.24    37.75             45.0       891.0      384.0        146.0        4.9
696    -122.10    37.69             41.0       746.0      387.0        161.0        3.9
```

```
In [119...
median = housing["total_bedrooms"].median()
sample_incomplete_rows["total_bedrooms"].fillna(median, inplace=True) # option 3
sample_incomplete_rows
```

```
Out[119...
      longitude  latitude  housing_median_age  total_rooms  total_bedrooms  population  househo
290    -122.16    37.77             47.0      1256.0           435.0      570.0        21
341    -122.17    37.75             38.0       992.0           435.0      732.0        25
538    -122.28    37.78             29.0     5154.0           435.0     3741.0       127
563    -122.24    37.75             45.0       891.0           435.0      384.0        14
696    -122.10    37.69             41.0       746.0           435.0      387.0        16
```

Could you think of another plausible imputation for this dataset? (Not graded)

Prepare Data

Recall we are trying to predict the median house value, our features will contain longitude, latitude, housing_median_age... and our target will be median_house_value

```
In [120...
housing_features = housing.drop("median_house_value", axis=1) # drop labels for
                                                             # the input to the model
housing_labels = housing["median_house_value"].copy()
```

```
In [121...
housing_features.head()
```

```
Out[121...
      longitude  latitude  housing_median_age  total_rooms  total_bedrooms  population  households
0    -122.23    37.88             41.0       880.0           129.0      322.0       126.0
1    -122.22    37.86             21.0     7099.0          1106.0     2401.0      1138.0
2    -122.24    37.85             52.0     1467.0           190.0      496.0       177.0
3    -122.25    37.85             52.0     1274.0           235.0      558.0       219.0
4    -122.25    37.85             52.0     1627.0           280.0      565.0       259.0
```

```
In [122...
```

```

# This cell implements the complete pipeline for preparing the data
# using sklearn's TransformerMixins
# Earlier we mentioned different types of features: categorical, and floats.
# In the case of floats we might want to convert them to categories.
# On the other hand categories in which are not already represented as integers
# feeding to the model.

# Additionally, categorical values could either be represented as one-hot vector
# Here we encode them using one hot vectors.

# DO NOT WORRY IF YOU DO NOT UNDERSTAND ALL THE STEPS OF THIS PIPELINE. CONCEPTS
# ONE-HOT ENCODING ETC. WILL ALL BE COVERED IN DISCUSSION

from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer

from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import OneHotEncoder

from sklearn.base import BaseEstimator, TransformerMixin

imputer = SimpleImputer(strategy="median") # use median imputation for missing v
housing_num = housing_features.drop("ocean_proximity", axis=1) # remove the cate
# column index
rooms_idx, bedrooms_idx, population_idx, households_idx = 3, 4, 5, 6

#
class AugmentFeatures(BaseEstimator, TransformerMixin):
    """
    implements the previous features we had defined
    housing["rooms_per_household"] = housing["total_rooms"]/housing["households"]
    housing["bedrooms_per_room"] = housing["total_bedrooms"]/housing["total_rooms"]
    housing["population_per_household"] = housing["population"]/housing["households"]
    """
    def __init__(self, add_bedrooms_per_room = True):
        self.add_bedrooms_per_room = add_bedrooms_per_room

    def fit(self, X, y=None):
        return self # nothing else to do

    def transform(self, X):
        rooms_per_household = X[:, rooms_idx] / X[:, households_idx]
        population_per_household = X[:, population_idx] / X[:, households_idx]
        if self.add_bedrooms_per_room:
            bedrooms_per_room = X[:, bedrooms_idx] / X[:, rooms_idx]
            return np.c_[X, rooms_per_household, population_per_household,
                          bedrooms_per_room]
        else:
            return np.c_[X, rooms_per_household, population_per_household]

attr_adder = AugmentFeatures(add_bedrooms_per_room=False)
housing_extra_attribs = attr_adder.transform(housing.values) # generate new features

# this will be a numerical pipeline
# 1. impute, 2. augment the feature set 3. normalize using StandardScaler()
num_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy="median")),
    ('attribs_adder', AugmentFeatures()),
    ('std_scaler', StandardScaler()),

```



```

    ])

housing_num_tr = num_pipeline.fit_transform(housing_num)

numerical_features = list(housing_num)
categorical_features = ["ocean_proximity"]

full_pipeline = ColumnTransformer([
    ("num", num_pipeline, numerical_features),
    ("cat", OneHotEncoder(), categorical_features),
])

housing_prepared = full_pipeline.fit_transform(housing_features)

```

Splitting our dataset

First we need to carve out our dataset into a training and testing cohort. To do this we'll use `train_test_split`, a very elementary tool that arbitrarily splits the data into training and testing cohorts.

```

In [123... from sklearn.model_selection import train_test_split
data_target = housing['median_house_value']
train, test, target, target_test = train_test_split(housing_prepared, data_target

```

Select a model and train

Once we have prepared the dataset it's time to choose a model.

As our task is to predict the `median_house_value` (a floating value), regression is well suited for this.

```

In [124... from sklearn.linear_model import LinearRegression

lin_reg = LinearRegression()
lin_reg.fit(train, target)

# let's try the full preprocessing pipeline on a few training instances
data = test
labels = target_test

print("Predictions:", lin_reg.predict(data)[:5])
print("Actual labels:", list(labels)[:5])

Predictions: [207828.06448011 281099.80175494 176021.36890539 93643.46744928
304674.47047758]
Actual labels: [136900.0, 241300.0, 200700.0, 72500.0, 460000.0]

```

```

In [125... from sklearn.metrics import mean_squared_error

preds = lin_reg.predict(test)
mse = mean_squared_error(target_test, preds)
rmse = np.sqrt(mse)
rmse

```


Out[125... 67879.86844243006

TODO: Applying the end-end ML steps to a different dataset.

We will apply what we've learnt to another dataset (airbnb dataset). We will predict airbnb price based on other features.

[35 pts] Visualizing Data

[5 pts] Load the data + statistics

- load the dataset
- display the first few rows of the data

In [126...

```
AIR_PATH = os.path.join("datasets", "airbnb")
def load_air_data(air_path):
    """
    loads AB_NYC_2019.csv dataset stored

    """
    csv_path = os.path.join(air_path, "AB_NYC_2019.csv")
    return pd.read_csv(csv_path)
airbnb = load_air_data(AIR_PATH)
airbnb.head()
```

Out[126...

	id	name	host_id	host_name	neighbourhood_group	neighbourhood	latitude
0	2539	Clean & quiet apt home by the park	2787	John	Brooklyn	Kensington	40.64749
1	2595	Skylit Midtown Castle	2845	Jennifer	Manhattan	Midtown	40.75362
2	3647	THE VILLAGE OF HARLEM....NEW YORK !	4632	Elisabeth	Manhattan	Harlem	40.80902
3	3831	Cozy Entire Floor of Brownstone	4869	LisaRoxanne	Brooklyn	Clinton Hill	40.68514
4	5022	Entire Apt: Spacious Studio/Loft by central park	7192	Laura	Manhattan	East Harlem	40.79851

- pull up info on the data type for each of the data fields. Will any of these be problematic feeding into your model (you may need to do a little research on this)? Discuss:

In [127...

```
airbnb.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48895 entries, 0 to 48894
Data columns (total 16 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     48895 non-null  int64
1   name                                  48879 non-null  object
2   host_id                               48895 non-null  int64
3   host_name                             48874 non-null  object
4   neighbourhood_group                   48895 non-null  object
5   neighbourhood                         48895 non-null  object
6   latitude                             48895 non-null  float64
7   longitude                             48895 non-null  float64
8   room_type                             48895 non-null  object
9   price                                 48895 non-null  int64
10  minimum_nights                        48895 non-null  int64
11  number_of_reviews                     48895 non-null  int64
12  last_review                           38843 non-null  object
13  reviews_per_month                     38843 non-null  float64
14  calculated_host_listings_count        48895 non-null  int64
15  availability_365                       48895 non-null  int64
dtypes: float64(3), int64(7), object(6)
memory usage: 6.0+ MB

```

[Response here] I think the problematic feeding into the model is the object Dtype. It might need to be converted into other primitive types before directly feed into the medel.

- drop the following columns: name, host_id, host_name, and last_review
- display a summary of the statistics of the loaded data

```
In [128... airbnb = airbnb.drop(["name", "host_id", "host_name", "last_review"], axis=1)
```

```
In [129... airbnb.describe()
```

```
Out[129...

```

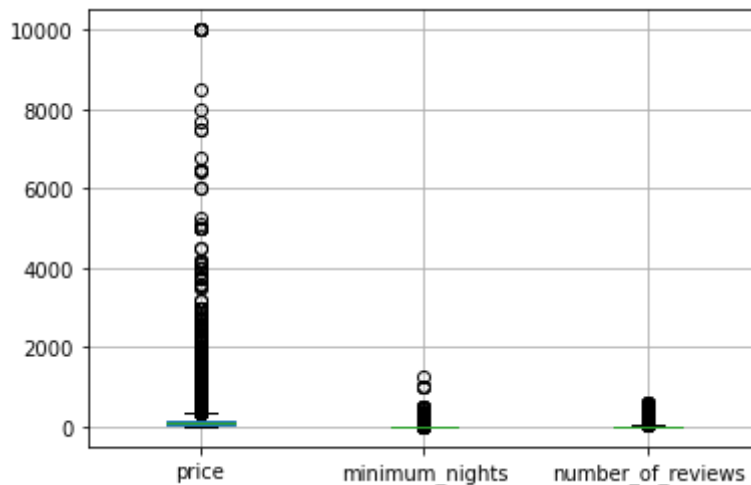
	id	latitude	longitude	price	minimum_nights	number_of_
count	4.889500e+04	48895.000000	48895.000000	48895.000000	48895.000000	48895.
mean	1.901714e+07	40.728949	-73.952170	152.720687	7.029962	23.
std	1.098311e+07	0.054530	0.046157	240.154170	20.510550	44.
min	2.539000e+03	40.499790	-74.244420	0.000000	1.000000	0.
25%	9.471945e+06	40.690100	-73.983070	69.000000	1.000000	1.
50%	1.967728e+07	40.723070	-73.955680	106.000000	3.000000	5.
75%	2.915218e+07	40.763115	-73.936275	175.000000	5.000000	24.
max	3.648724e+07	40.913060	-73.712990	10000.000000	1250.000000	629.

[5 pts] Boxplot 3 features of your choice

- plot boxplots for 3 features of your choice

```
In [130... airbnb.boxplot(['price', 'minimum_nights', 'number_of_reviews'])
```

Out[130... <matplotlib.axes._subplots.AxesSubplot at 0x7fb397f78fa0>



- describe what you expected to see with these features and what you actually observed


[Response here] Since the boxplot is a standardized way of displaying the distribution of data based on "minimum", first quartile (Q1), median, third quartile (Q3), and "maximum". I expected it to tell me about the outliers and values of the three features "price", "min_nights", and "number_of_reviews" I chose. However, according to the result generated from the plot, all the points are closely printed which is difficult for me to distinguish them. I think it might be because the data is not evenly distributed.

High variability in price with long tail values, review numbers much more compact, however availability has a wider variance.

[10 pts] Plot average price of a listing per neighbourhood_group

```
In [131... meanGrouped = airbnb.groupby(['neighbourhood_group']).mean()['price']
meanGrouped.plot(kind='bar', rot=0)
plt.ylabel("Price", fontsize=12)
plt.xlabel("Neighbourhood Group", fontsize=12)
plt.title("Average Price of a Listing per Neighbourhood Group", fontsize=12)
```

Out[131... Text(0.5, 1.0, 'Average Price of a Listing per Neighbourhood Group')



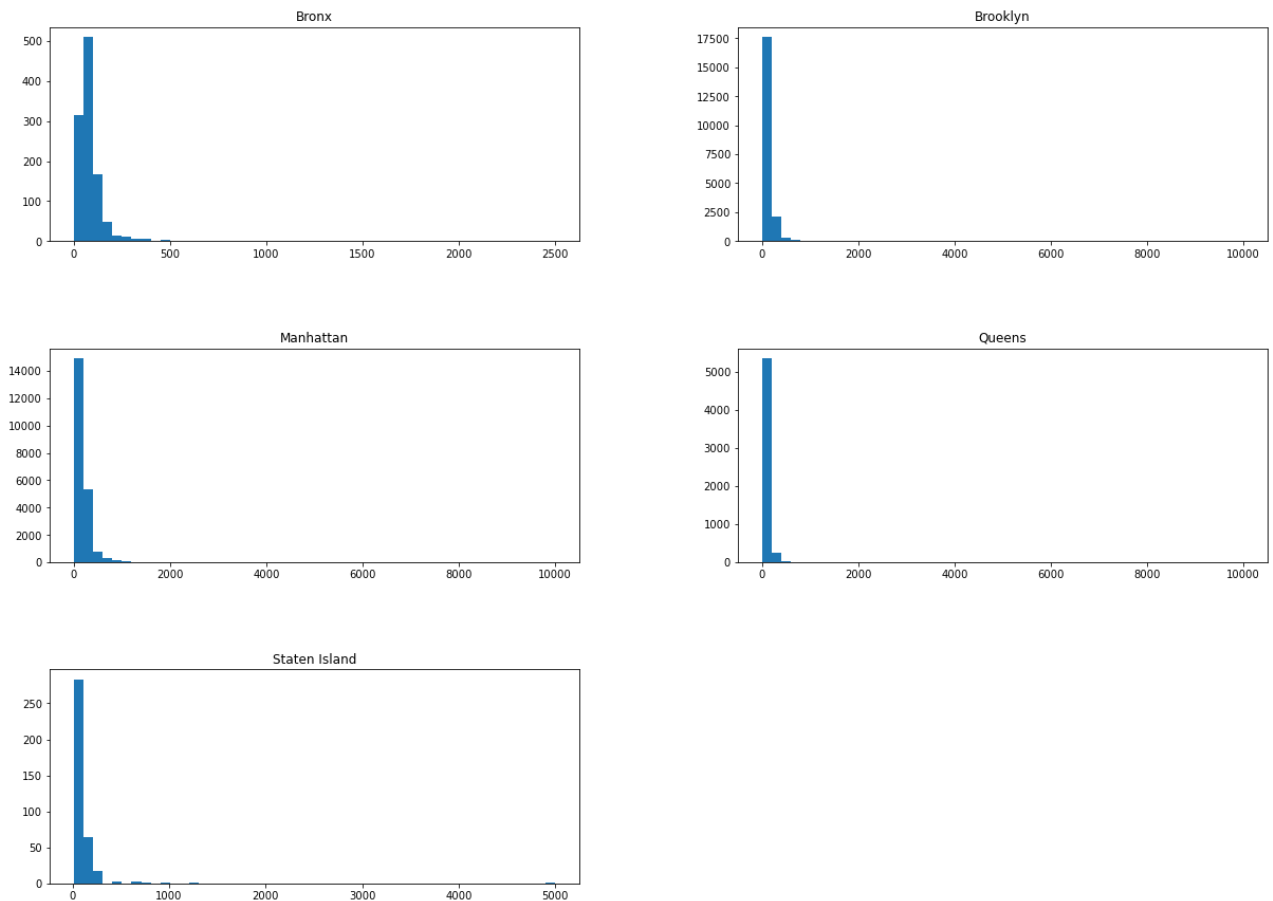
- describe what you expected to see with these features and what you actually observed

[Response here] What I expected here is that the average price of the listings are different depending on the varies Neighborhoo group, since the most important element for gaining more is the location of the houses. According to the actually generated plot, we can see the price are in the range from about 75 to 200 because of their different neighbourhood group, and it's clearly that Manhattan's average price is higher than other areas such as Bronx.

- So we can see different neighborhoods have dramatically different pricepoints, but how does the price breakdown by range. To see let's do a histogram of price by neighborhood to get a better sense of the distribution.

```
In [132... airbnb['price'].hist(by = airbnb['neighbourhood_group'], bins=50, figsize=(20,15
```

```
Out[132... array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7fb38efb4370>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7fb38ef8d490>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x7fb38ed6cc10>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7fb38bbef3d0>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x7fb38f451b50>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7fb38f458370>]],
dtype=object)
```



[5 pts] Plot map of airbnbs throughout New York (if it gets too crowded take a subset of the data, and try to make it look nice if you can :)).

In [133...


```
# load an image of new york
images_path = os.path.join('.', "images")
os.makedirs(images_path, exist_ok=True)
filename = "newyork.png"
import matplotlib.image as mpimg
newYork_img=mpimg.imread(os.path.join(images_path, filename))
ax = airbnb.plot(kind="scatter", x="longitude", y="latitude", figsize=(10,7),
                  s=airbnb['price']/100, label="price",
                  c="price", cmap=plt.get_cmap("jet"),
                  colorbar=False, alpha=0.4,
                  )

# overlay the new york map on the plotted scatter plot
# note: plt.imshow still refers to the most recent figure
# that hasn't been plotted yet.
plt.imshow(newYork_img, extent=[-74.24, -73.71, 40.488, 40.91], alpha=0.5,
           cmap=plt.get_cmap("jet"))
plt.ylabel("Latitude", fontsize=14)
plt.xlabel("Longitude", fontsize=14)

# setting up heatmap colors based on price feature
prices = airbnb["price"]
tick_values = np.linspace(prices.min(), prices.max(), 11)
cb = plt.colorbar()
cb.ax.set_yticklabels(["$%dk"%(round(v/1000)) for v in tick_values], fontsize=14)
cb.set_label('Price', fontsize=16)
```

```
plt.legend(fontsize=16)
save_fig("ny_housing_prices_plot")
plt.show()
```

Saving figure ny_housing_prices_plot




[5 pts] Plot average price of room types who have availability greater than 180 days and neighbourhood_group is Manhattan

```
In [134... avgPriceGt180InMht = airbnb[(airbnb['availability_365'] > 180) & (airbnb['neighb
avgPriceGt180InMht.plot(kind='bar', rot= 0)
plt.ylabel("Average Price")
plt.xlabel("Room Types")
plt.title("Average Price of Room Types who have Availability Greater than 180 Da
```

```
Out[134... Text(0.5, 1.0, 'Average Price of Room Types who have Availability Greater than 1
80 Days in Manhattan')
```

Average Price of Room Types who have Availability Greater than 180 Days in Manhattan



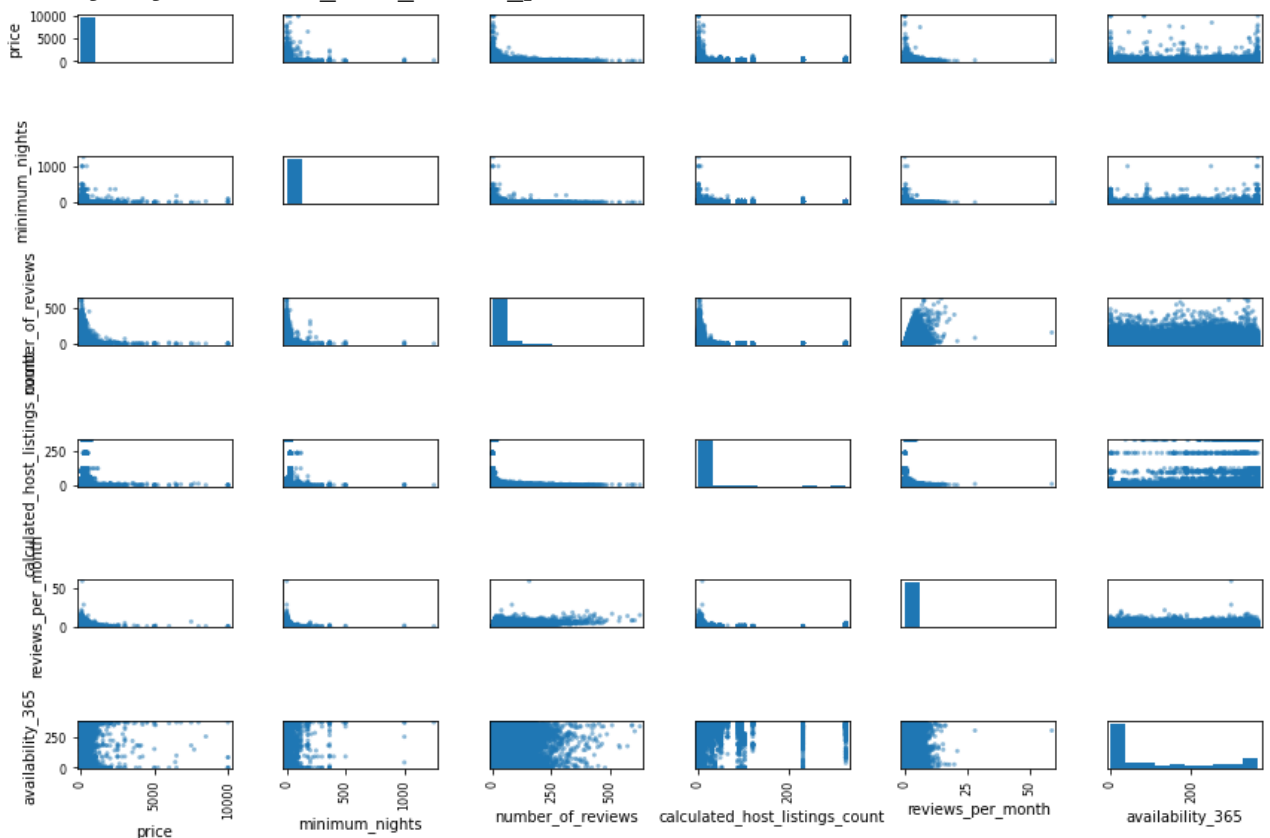
[5 pts] Plot correlation matrix

- which features have positive correlation?
- which features have negative correlation?

In [135...

```
# airbnb.corr()
# price minimum_nights number_of_reviews reviews_per_month calculated_host_listings_count
features = ['price', 'minimum_nights', 'number_of_reviews', 'calculated_host_listings_count', 'reviews_per_month', 'availability_365']
scatter_matrix(airbnb[features], figsize=(12, 8))
save_fig("Airbnb_corr_matrix_plot")
```

Saving figure Airbnb_corr_matrix_plot



[Response here] Reviews_per_month and number_of_reviews has positive correlation. Price

and minimum_nights has negative correlation. Price and number_of_reviews has negative correlation. Price and reviews_per_month has negative correlation.

[30 pts] Prepare the Data

[5 pts] Augment the dataframe with two other features which you think would be useful

```
In [136... # price minimum_nights number_of_reviews reviews_per_month calculate
airbnb['review_monthly_rate'] = airbnb['number_of_reviews'] / airbnb['reviews_per_month']
airbnb['book_status'] = airbnb['availability_365'] / airbnb['minimum_nights']
```

[5 pts] Impute any missing feature with a method of your choice, and briefly discuss why you chose this imputation method

```
In [137... airbnb_incomplete_rows = airbnb[airbnb.isnull().any(axis=1)].head()
airbnb_incomplete_rows
```

```
Out[137...      id  neighbourhood_group  neighbourhood  latitude  longitude  room_type  price  minimum_nights
2    3647                Manhattan           Harlem    40.80902   -73.94190   Private room    150
19   7750                Manhattan       East Harlem    40.79685   -73.94872   Entire home/apt    190
26   8700                Manhattan           Inwood    40.86754   -73.92639   Private room     80
36  11452                Brooklyn  Bedford-Stuyvesant    40.68876   -73.94312   Private room     35
38  11943                Brooklyn           Flatbush    40.63702   -73.96327   Private room    150
```

```
In [138... airbnb["reviews_per_month"].fillna(median, inplace=True)
airbnb["review_monthly_rate"].fillna(median, inplace=True)
# I choose to replace na values with median values instead of drop the feature or
# in order to avoid the missing data and preserve as more data as possible.
airbnb
```

```
Out[138...      id  neighbourhood_group  neighbourhood  latitude  longitude  room_type  price
0    2539                Brooklyn    Kensington    40.64749   -73.97237   Private room    149
1    2595                Manhattan           Midtown    40.75362   -73.98377   Entire home/apt    225
2    3647                Manhattan           Harlem    40.80902   -73.94190   Private room    150
3    3831                Brooklyn    Clinton Hill    40.68514   -73.95976   Entire home/apt     89
```


	id	neighbourhood_group	neighbourhood	latitude	longitude	room_type	price
4	5022	Manhattan	East Harlem	40.79851	-73.94399	Entire home/apt	80
...
48890	36484665	Brooklyn	Bedford-Stuyvesant	40.67853	-73.94995	Private room	70
48891	36485057	Brooklyn	Bushwick	40.70184	-73.93317	Private room	40
48892	36485431	Manhattan	Harlem	40.81475	-73.94867	Entire home/apt	115
48893	36485609	Manhattan	Hell's Kitchen	40.75751	-73.99112	Shared room	55
48894	36487245	Manhattan	Hell's Kitchen	40.76404	-73.98933	Private room	90

48895 rows × 14 columns

[15 pts] Code complete data pipeline using sklearn mixins

In [139...

```
airbnb_df = load_air_data(AIR_PATH).drop(
    ["name", "host_id", "host_name", "last_review", "id", "latitude", "longitude"]
)
imputer = SimpleImputer(strategy="median")
categorical_features = ["neighbourhood_group", "neighbourhood", "room_type"]
feature_drop = airbnb_df.drop(categorical_features, axis=1)
min_nights_id, num_reviews_id, review_per_month_id, availability_id, = 1, 2, 3, 4

class NewAugmentFeatures(BaseEstimator, TransformerMixin):
    def __init__(self):
        pass
    def fit(self, X, y=None):
        return self
    def transform(self, X):
        review_monthly_rate = X[:, num_reviews_id] / X[:, review_per_month_id]
        book_status = X[:, availability_id] / X[:, min_nights_id]
        return np.c_[X, review_monthly_rate, book_status]

num_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy="median")),
    ('attrs_adder', NewAugmentFeatures()),
    ('std_scaler', StandardScaler()),
])

full_pipeline = ColumnTransformer([
    ("num", num_pipeline, list(feature_drop)),
    ("cat", OneHotEncoder(), categorical_features),
])

air_drop = load_air_data(AIR_PATH).drop(
    ["id", "latitude", "longitude", "name", "host_id", "host_name", "last_review"]
)
airbnb_prepared = full_pipeline.fit_transform(air_drop)
xtest = air_drop['price']
```

[5 pts] Set aside 20% of the data as test test (80% train, 20% test).

```
In [140... from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(airbnb_prepared, xtest, test
```

[15 pts] Fit a model of your choice

The task is to predict the price, you could refer to the housing example on how to train and evaluate your model using MSE. Provide both test and train set MSE values.

```
In [142... from sklearn.linear_model import LinearRegression
lin_reg = LinearRegression()
lin_reg.fit(x_train, y_train)

xpreds= lin_reg.predict(x_test)
print("Predict result: ", xpreds[:5])
print("Actual result:", list(y_test)[:5])
mse = mean_squared_error(y_test, xpreds)
print("Test MSE: ", mse)

ypreds= lin_reg.predict(x_train)
tmse = mean_squared_error(y_train, ypreds)
print("Train MSE: ", tmse)
```

```
Predict result: [225.00000194 649.00000525 299.99999951 26.00000882 125.000000
01]
Actual result: [225, 649, 300, 26, 125]
Test MSE: 3.1165441874167874e-10
Train MSE: 5.876895228476491e-10
```

```
In [ ]:
```